



Core idea: convert EEG into retrievable semantic evidence, then let the LLM generate a conservative caption from reliable cues.